

MDP: Markov Decision Problem



Outline

- **Markov decision problem**
- Optimal state-value function and optimal action-value function
- Value iteration
- Example



Markov Decision Problem (MDP)

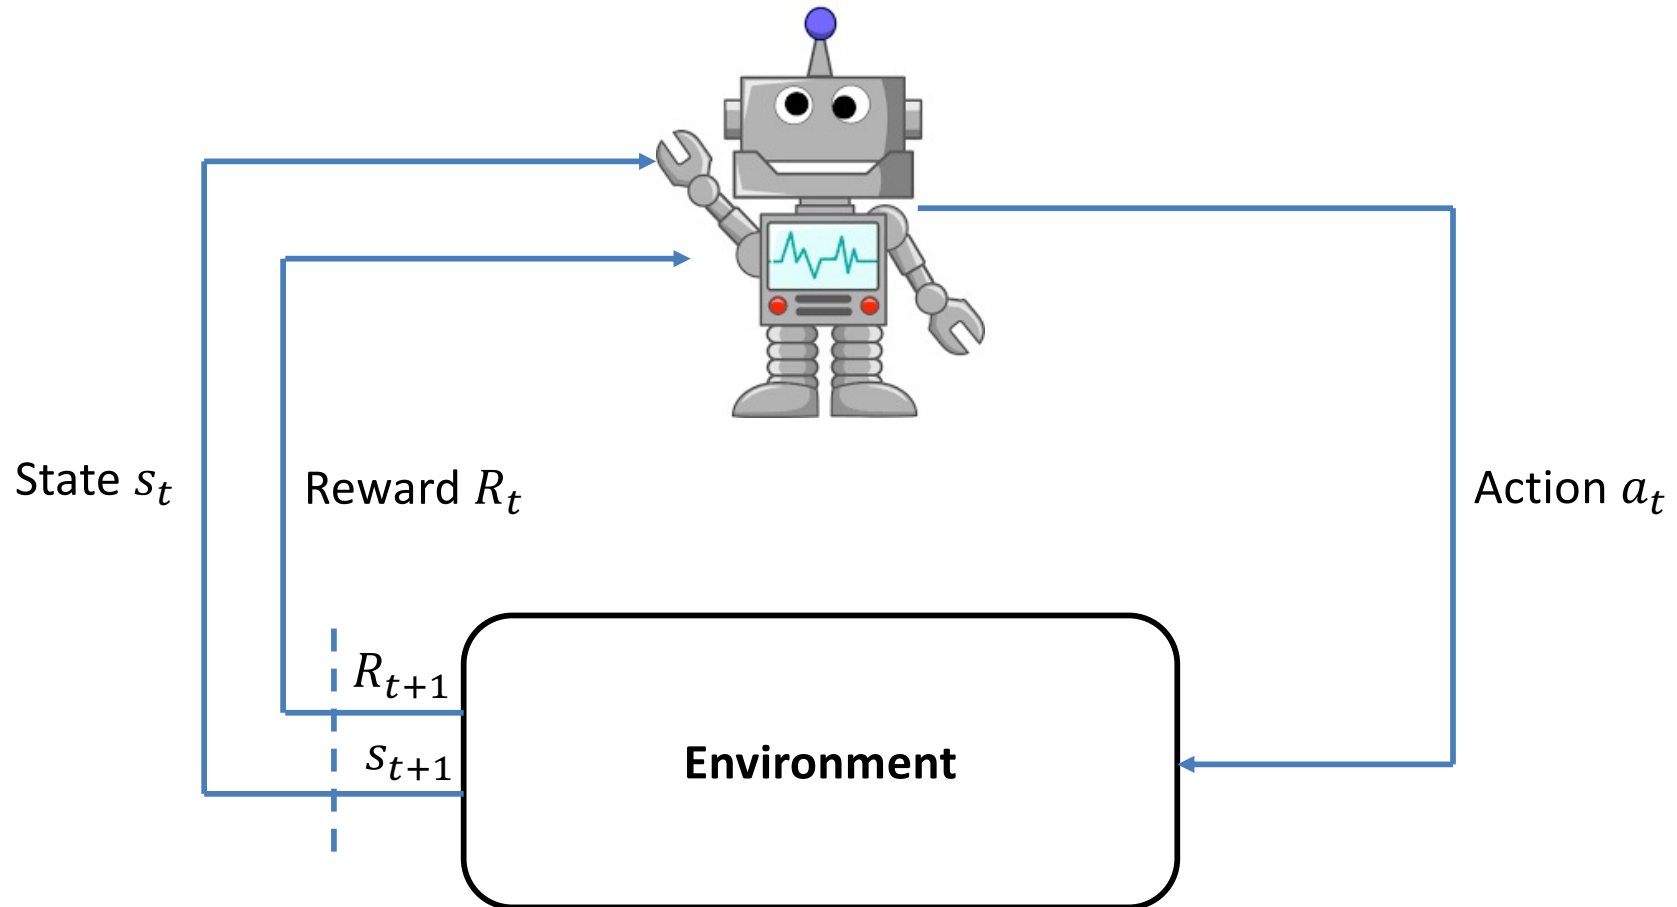
- A framework for **sequential decision making** of a single agent
- **Markovian transition model** and **fully observable**
 - i.e., we assume it verifies the Markov property and $s_t = \theta_t$
- **Planning horizon** can be infinite



Markov Decision Problem (MDP)

- We can formally define an MDP with following elements:
 - **Discrete time** $t = 0, 1, 2, \dots$
 - **A discrete set of states** $s \in S$
 - **A discrete set of actions** $a \in A$
 - **A stochastic transition model** $P(s'|s, a)$
 - the world transitions stochastically to state s' when the agent takes action a at state s
 - **A reward function** $R: S \times A \rightarrow \mathbb{R}$
 - An agent receives a reward $R(s, a)$ when it takes action a at state s
 - **A discount rate** γ

Markov Decision Problem (MDP)




Markov Decision Problem (MDP)

- The **task** of the agent is to **maximize a function of accumulated reward** over its planning horizon
- A standard function is the **discounted future reward**:

$$R(s_t, a_t) + \gamma R(s_{t+1}, a_{t+1}) + \gamma^2 R(s_{t+2}, a_{t+2}) + \dots$$

- where γ is the **discount rate** and $0 \leq \gamma \leq 1$
 - γ is a value between 0 and 1 in order to ensure the sum remains finite for infinite horizon

Which trajectory would you choose?



What is the best sequence of actions in this MDP?

- Discounted future reward for the blue trajectory:

$$0 + \gamma 0 + \gamma^2 0 + \gamma^3 0 + \gamma^4 0 + \gamma^5 0 + \gamma^6 0 + \gamma^7 0 + \gamma^8 100$$

- Discounted future reward for the red trajectory:

$$0 + \gamma 0 + \gamma^2 100$$

- The red sequence has the highest value!

Markov Decision Problem (MDP)

- A **stationary policy** $\pi : S \rightarrow A$ of the agent in an MDP is a mapping $\pi(s)$ from states to actions.

$$\pi: S \rightarrow A$$

- **Different policies** will produce **different discounted future rewards**
 - Each policy will take the agent through different trajectories

Markov Decision Problem (MDP)

- **Definition:** the **state-value function** of a state s under a policy π is the expected return the agent can receive when starting in state s and then following policy π :

$$V^{\pi}(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, a_t = \pi(s_t) \right]$$

Definition: the **action-value function (Q-values)** of taking an action a in state s under a policy π is the expected return the agent can receive when starting in state s , taking action a , and then following policy π :

$$Q^{\pi}(s, a) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, a_0 = a, a_{t>0} = \pi(s_t) \right]$$

Outline

- Markov decision problem
- **Optimal state-value function and optimal action-value function**
- Value iteration
- Example



Markov Decision Problem (MDP)

- A **policy π** is defined to be **better than or equal to a policy π'** if the expected return of π is greater or equal to expected return of π' :

$$\pi \geq \pi' \text{ if and only if } V^\pi(s) \geq V^{\pi'}(s), \text{ for all } s \in S$$

- **There is always at least one policy** that is better than or equal to all other policies, which is the **optimal policy**
 - Note that an MDP might have more than one optimal policy
 - We denote all the optimal policies by π^*

Markov Decision Problem (MDP)

- These policies π^* share the same state-value function, called **optimal state-value function**, with the following definition:

$$V^*(s) = \max_{\pi} V^{\pi}(s), \text{ for all } s \in S$$

- These policies π^* also share the same action-value function, called **optimal action-value function**, with the following definition:

$$Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a), \text{ for all } s \in S \text{ and } a \in A$$

Markov Decision Problem (MDP)

- But how do we find this optimal policy?
 - Naïve solution: compute the state-value function for all policies

Outline

- Markov decision problem
- Optimal state-value function and optimal action-value function
- **Value iteration**
- Example



Markov Decision Problem (MDP)

- Lets us now define the **stationary policy** π of the agent in an MDP as a mapping from states to a distribution over actions

$$\pi: S \rightarrow \Delta(A)$$

- An example of a distribution for an MDP with $S = \{x, y\}$, $A = \{b, c\}$, P, R :

$$\pi = \begin{matrix} & \begin{matrix} \textcolor{red}{b} \\ \textcolor{red}{\downarrow} \end{matrix} & \begin{matrix} \textcolor{red}{c} \\ \textcolor{red}{\downarrow} \end{matrix} \\ \begin{matrix} \textcolor{red}{x} \\ \textcolor{red}{\downarrow} \end{matrix} & \begin{bmatrix} 1 & 0 \\ 0.5 & 0.5 \end{bmatrix} \\ \begin{matrix} \textcolor{red}{y} \\ \textcolor{red}{\downarrow} \end{matrix} & \end{matrix}$$

- Thus, we can use the notation $\pi(a|s)$
 - For instance, $\pi(a = b|s = x) = 1$

Markov Decision Problem (MDP)

$$V^{\pi}(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \mid s_0 = s, a_t = \pi(s_t) \right]$$

$$V^{\pi}(s) = E_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t R_t \mid s_0 = s \right]$$

$$V^{\pi}(s) = E_{\pi} [R_0 + \gamma R_1 + \gamma^2 R_2 + \cdots \mid s_0 = s]$$

Markov Decision Problem (MDP)

$$V^\pi(s) = E_\pi[R_0 + \gamma R_1 + \gamma^2 R_2 + \cdots | s_0 = s]$$

$$V^\pi(s) = E_\pi[R_0 + \gamma(R_1 + \gamma R_2 + \cdots) | s_0 = s]$$

$$V^\pi(s) = E_\pi[R_0 | s_0 = s] + E_\pi[\gamma(R_1 + \gamma R_2 + \cdots) | s_0 = s]$$

$$V^\pi(s) = R_\pi(s) + E_\pi[\gamma(R_1 + \gamma R_2 + \cdots) | s_0 = s]$$



$$R_\pi(s) = E_\pi[R(s, a)] = \sum_{a \in A} \pi(a|s) R(s, a)$$

Policy-averaged reward

Markov Decision Problem (MDP)

We would like this
to be s_1

$$V^\pi(s) = R_\pi(s) + \gamma E_\pi[(R_1 + \gamma R_2 + \dots) | s_0 = s]$$

We use the law of total
probability with expectation

$$V^\pi(s) = R_\pi(s) + \gamma \sum_{s' \in S} E_\pi[(R_1 + \gamma R_2 + \dots) | s_1 = s'] P_\pi(s' | s)$$

$$\begin{aligned} P_\pi(s' | s) &= E_\pi[P(s' | s, a)] \\ &= \sum_{a \in A} \pi(a | s) P(s' | s, a) \end{aligned}$$

Policy-averaged
probabilities

$$V^\pi(s) = R_\pi(s) + \gamma \sum_{s' \in S} P_\pi(s' | s) V^\pi(s')$$

Markov Decision Problem (MDP)

- We now take recursion

$$V^\pi(s) = R_\pi(s) + \gamma \sum_{s' \in S} P_\pi(s'|s) V^\pi(s')$$

- And make the following substitutions

$$V^\pi(s) = \sum_{a \in A} \pi(a|s) R(s, a) + \gamma \sum_{s' \in S} \sum_{a \in A} \pi(a|s) P(s'|s, a) V^\pi(s')$$

$$V^\pi(s) = \sum_{a \in A} \pi(a|s) \left[\underbrace{R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^\pi(s')}_{Q^\pi(s, a)} \right]$$

Markov Decision Problem (MDP)

- The **optimal state-value function** can be written in a special form without reference to a specific policy (so-called **Bellman equation**):

$$V^*(s) = \max_{a \in A} \left[R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^*(s') \right]$$

- A similar recursive equation holds for the **Q-values**:

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \max_{a' \in A} Q^*(s', a')$$

Value Iteration

- **Value iteration** is a simple and efficient method for computing optimal values in an MDP

Value Iteration

- We initialize arbitrarily a state-value function (e.g., with zeros, ones, etc.)
- Then we iteratively apply the Bellman equation (in the previous slides) turned into an assignment operation:

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

- Repeat the above two equations until V does not change significantly between two consecutive steps

Value Iteration

- Value iteration converges to the optimal Q^* for any initialization
- After computing the optimal Q^* , we can extract the policy as follows:

$$\pi^*(s) \in \operatorname{argmax}_{a \in A} Q^*(s, a)$$

Outline

- Markov decision problem
- Optimal state-value function and optimal action-value function
- Value iteration
- **Example**



Example


- We have the following MDP:

- $S = \{1, 2, 3\}$

- $A = \{left, right\}$

- $P(s'|s, a = left) = \begin{bmatrix} 1 & 0 & 0 \\ 0.8 & 0.2 & 0 \\ 0 & 0.8 & 0.2 \end{bmatrix}$

- $P(s'|s, a = right) = \begin{bmatrix} 0.2 & 0.8 & 0 \\ 0 & 0.2 & 0.8 \\ 0 & 0 & 1 \end{bmatrix}$




Example

- We have the following MDP:

- $R(s, a) = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 1 & 1 \end{bmatrix}$

- $\gamma = 0.9$



Example

- Let us use value iteration:
 - We initialize state-value function with ones

$$V = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

Example

- we iteratively apply the following equation ($i = 0$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q(1, left) := 0 + 0.9(1 \times 1 + 0 \times 1 + 0 \times 1) = 0.9$$

$$Q(2, left) := 0 + 0.9(0.8 \times 1 + 0.2 \times 1 + 0 \times 1) = 0.9$$

$$Q(3, left) := 1 + 0.9(0 \times 1 + 0.8 \times 1 + 0.2 \times 1) = 1.9$$

$$Q(1, right) := 0 + 0.9(0.2 \times 1 + 0.8 \times 1 + 0 \times 1) = 0.9$$

$$Q(2, right) := 0 + 0.9(0 \times 1 + 0.2 \times 1 + 0.8 \times 1) = 0.9$$

$$Q(3, right) := 1 + 0.9(0 \times 1 + 0 \times 1 + 1 \times 1) = 1.9$$

Example

- we iteratively apply the following equations ($i = 0$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q = \begin{bmatrix} 0.9 & 0.9 \\ 0.9 & 0.9 \\ 1.9 & 1.9 \end{bmatrix}$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

$$V = \begin{bmatrix} 0.9 \\ 0.9 \\ 1.9 \end{bmatrix}$$

Example

- we iteratively apply the following equation ($i = 1$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q(1, left) := 0 + 0.9(1 \times 0.9 + 0 \times 0.9 + 0 \times 1.9) = 0.81$$

$$Q(2, left) := 0 + 0.9(0.8 \times 0.9 + 0.2 \times 0.9 + 0 \times 1.9) = 0.81$$

$$Q(3, left) := 1 + 0.9(0 \times 0.9 + 0.8 \times 0.9 + 0.2 \times 1.9) = 1.99$$

$$Q(1, right) := 0 + 0.9(0.2 \times 0.9 + 0.8 \times 0.9 + 0 \times 1.9) = 0.81$$

$$Q(2, right) := 0 + 0.9(0 \times 0.9 + 0.2 \times 0.9 + 0.8 \times 1.9) = 1.53$$

$$Q(3, right) := 1 + 0.9(0 \times 0.9 + 0 \times 0.9 + 1 \times 1.9) = 2.71$$

Example

- we iteratively apply the following equations ($i = 1$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q = \begin{bmatrix} 0.81 & 0.81 \\ 0.81 & 1.53 \\ 1.99 & 2.71 \end{bmatrix}$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

$$V = \begin{bmatrix} 0.81 \\ 1.53 \\ 2.71 \end{bmatrix}$$

Example

- we iteratively apply the following equations ($i = 2$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q = \begin{bmatrix} 0.73 & 1.25 \\ 0.86 & 2.23 \\ 2.59 & 3.44 \end{bmatrix}$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

$$V = \begin{bmatrix} 1.25 \\ 2.23 \\ 3.44 \end{bmatrix}$$

Example

- After a few iterations...

Example

- we iteratively apply the following equations ($i = 47$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q = \begin{bmatrix} 6.88 & 7.65 \\ 7.07 & 8.72 \\ 9.06 & 9.94 \end{bmatrix}$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

$$V = \begin{bmatrix} 7.65 \\ 8.72 \\ 9.94 \end{bmatrix}$$

Example

- we iteratively apply the following equations ($i = 48$)

$$Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s'), \quad \forall s, \forall a$$

$$Q = \begin{bmatrix} 6.89 & 7.66 \\ 7.08 & 8.73 \\ 9.07 & 9.95 \end{bmatrix}$$

$$V(s) := \max_{a \in A} Q(s, a), \quad \forall s$$

$$V = \begin{bmatrix} 7.66 \\ 8.73 \\ 9.95 \end{bmatrix}$$

Example

- The algorithm repeated until V did not change significantly between two consecutive steps.
- Result:

$$V^* = \begin{bmatrix} 7.66 \\ 8.73 \\ 9.95 \end{bmatrix}$$

$$Q^* = \begin{bmatrix} 6.89 & 7.66 \\ 7.08 & 8.73 \\ 9.07 & 9.95 \end{bmatrix}$$


Example

- After computing the optimal Q^* , we can extract the policy as follows:

$$\pi^*(s) \in \operatorname{argmax}_{a \in A} Q^*(s, a)$$

$$\pi^* = \begin{bmatrix} \textit{right} \\ \textit{right} \\ \textit{right} \end{bmatrix}$$

$$\pi^* = \begin{bmatrix} 0 & 1 \\ 0 & 1 \\ 0 & 1 \end{bmatrix}$$



Python Code

```
import numpy as np
np.set_printoptions(precision=2, suppress=True)

# States
S = ['1', '2', '3']

# Actions
A = ['L', 'R']

# Transition probabilities

L = np.array([[1.0, 0.0, 0.0],
              [0.8, 0.2, 0.0],
              [0.0, 0.8, 0.2]])

R = np.array([[0.2, 0.8, 0.0],
              [0.0, 0.2, 0.8],
              [0.0, 0.0, 1.0]])

P = [L, R]

# Reward function

R = np.array([[0.0, 0.0],
              [0.0, 0.0],
              [1.0, 1.0]])

gamma = 0.9
```

Python Code

```
# Initialize V
V = np.ones(len(S))

err = 1
i = 0

while err > 1e-2:

    Q = []

    # Compute Q-values associated with each action
    for a in range(len(A)):
        Q += [R[:, a] + gamma * P[a].dot(V)]

    print('Q values at time ', i)
    print(Q)

    # Compute maximum for each state
    Vnew = np.max(Q, axis=0)
    print('V-function at time ', i)
    print(Vnew)

    # Compute error
    err = np.linalg.norm(V - Vnew)

    # Update V
    V = Vnew

    i += 1

print('\nV* =')
print(V[:, None])

print('\nQ* =')
print(Q)
```


Thank You



sardinha@inf.puc-rio.br